



Course: **EE321 Intro to AI**

Complex Engineering Problem

Name: M RAYAN UR REHMAN KHAN

Reg.: 2023281

IMPORTANT: This is an INDIVIDUAL project

You must run the code yourself.

The exam will ask specific questions about YOUR model's behavior.
If you copy from the internet, you will likely to fail the viva/exam.

Your Must

Build an appropriate AI model that can tell whether a picture contains a **hot dog** or **not a hot dog**.

Learning Objectives

By the end of this project, you will be able to:

- Load and explore image data
- Build an ANN and a CNN in TensorFlow/Keras
- Identify and fix overfitting using Dropout and Data Augmentation
- Plot and interpret training curves
- Test your model on real-world images

Dataset

You will use the **Hot Dog / Not Hot Dog** dataset from Kaggle:
<https://www.kaggle.com/dansbecker/hot-dog-not-hot-dog>

What You Need

- A Google account (for Google Colab)
- The dataset downloaded to your computer (upload to Colab)
- **Submit your overleaf generated PDF by May 07, 2026 by 12:00 PM**

Grading Rubric

Task	Points
Step 2: Data Exploration (answers)	5
Step 3: Visualization (screenshot)	5
Step 5: ANN Model (accuracy + analysis)	10
Step 6: CNN Model (accuracy + comparison)	15
Step 7: Overfitting Fix (before/after plots)	20
Step 8: Real Image Test (screenshot)	10
Step 9: Answers to all questions	20
Step 10: YOLO Extension (extra credit)	5
Total	85 + 5 EC

Step 1: Setup Google Colab

Follow these instructions carefully:

How to Run This Project in Google Colab

1. Goto <https://colab.research.google.com>
2. Click **File** → **New Notebook**
3. Rename it: **HotDog Classifier YourName.ipynb**
4. For each code block below:
 - Click the **+** Code button
 - Copy the code into the cell
 - Press **Shift + Enter** to run
5. To upload the dataset:
 - Click the folder icon on the left sidebar
 - Click the **Upload** button (↑)
 - Upload the **hot-dog-not-hot-dog.zip** file
 - Run: `!unzip hot-dog-not-hot-dog.zip`

Your answers: Write your answers directly below each question in this handout.

Step 2: Import Libraries

Run this code in Colab:

```
import numpy as np import matplotlib.pyplot as plt import tensorflow as
tf from tensorflow import keras from tensorflow.keras import layers from
tensorflow.keras.preprocessing.image import ImageDataGenerator import os import random

print("Libraries loaded successfully!") print("TensorFlow version:",
tf.__version__)
```

1. What TensorFlow version are you using?

Ans: I am using 2.20.0

Step 3: Load and Explore the Data

Run this code to see how many images you have:

```
# Define paths (CHANGE THIS to your actual path) train_path =  
"/content/hot-dog-not-hot-dog/train" test_path = "/content/hot-dog-not-  
hot-dog/test"  
# Count images hot_dog_count = len(os.listdir(os.path.join(train_path, "hot_dog"))) not_hot_dog_count  
= len(os.listdir(os.path.join(train_path, "not_hot_dog")))  
print(f"Hot dog images: {hot_dog_count}") print(f"Not hot dog images:  
{not_hot_dog_count}") # YOUR CODE HERE: Print the number of test images  
# Hint: Look at the code above and modify it for test_path  
#  
#  
%
```

1. How many training images are in each class?

Ans: 249 images for hot_dog and 249 images for not_hot_dog

2. Is the dataset balanced? Why does that matter?

Ans: the dataset is perfectly balanced. This matters because it ensures the model doesn't develop a bias toward the majority class which would lead to misleadingly high accuracy by simply guessing the more common label

Step 4: Visualize the Data

```
# Display 6 random images plt.figure(figsize=(12, 6)) for i in range(6): # Pick random class
class_name = random.choice(["hot_dog", "not_hot_dog"]) class_path =
os.path.join(train_path, class_name) img_name = random.choice(os.listdir(class_path))
img_path = os.path.join(class_path, img_name)
    # Load and show image img =
    plt.imread(img_path) plt.subplot(2, 3,
    i+1) plt.imshow(img)
    plt.title(class_name) plt.axis('off')
plt.tight_layout() plt.show()
# YOUR CODE HERE: Print the shape (height, width, channels) of one image. Hint: Use img.shape after
loading an image
#
#
```

1. Paste a screenshot of your 6 random images here:

Sample Training Images

Not Hot Dog



Hot Dog



Hot Dog



Not Hot Dog



Not Hot Dog



Not Hot Dog



2. What is the shape (height, width, channels) of a typical image? (Use img.shape)

Ans: 64, 64, 3

Step 5: Prepare the Data Generator

```
# Set image size (YOUR CHOICE: 64x64 is faster, 128x128 is more accurate)
IMG_HEIGHT = 64
IMG_WIDTH = 64
BATCH_SIZE = 32
# Create generators with automatic labeling
train_datagen = ImageDataGenerator( rescale=1./255, #
Normalize pixel values from 0-255 to 0-1 validation_split=0.2 # Use 20% of training data for validation
)

train_generator = train_datagen.flow_from_directory( train_path, target_size=(IMG_HEIGHT,
IMG_WIDTH), batch_size=BATCH_SIZE, class_mode='binary', # Binary classification (hot dog
vs not) subset='training' # Training set (80% of data)
)

validation_generator = train_datagen.flow_from_directory( train_path,
target_size=(IMG_HEIGHT, IMG_WIDTH), batch_size=BATCH_SIZE,
class_mode='binary', subset='validation' # Validation set (20% of data)
)

# YOUR CODE HERE: Print the class labels
#
# Hint: Print train_generator.class_indices
```

1. What does class label 0 represent? What does 1 represent

Ans: Label 0 represents hot_dog. Label 1 represents
Not_hot_dog

Step 6: Build a Simple ANN (It Will Fail)

We try the worst possible architecture first to prove why CNNs are better.

```
# Build a simple ANN ann_model = keras.Sequential([ # Flatten the 2D image into 1D
layers.Flatten(input_shape=(IMG_HEIGHT, IMG_WIDTH, 3)),
    # Dense layers layers.Dense(128, activation='relu'), layers.Dense(64, activation='relu'),
    layers.Dense(1, activation='sigmoid') # Output: probability of " hot dog"
])
# Compile the model ann_model.compile(optimizer='adam', loss='binary_crossentropy',
metrics=['accuracy']) # Show model summary ann_model.summary() # Train for 10 epochs
history_ann = ann_model.fit(train_generator, validation_data= validation_generator, epochs=10)
# YOUR CODE HERE: Plot training vs validation accuracy.
#
#
# Hint: Use
#         plt.plot(history_ann.history['XXXX'], label='XXXX')
#         plt.plot(history_ann.history['XXXX'], label='XXXX')
```



1. What was the final validation accuracy of the ANN?

Ans: 57.14%

2. Why does the ANN perform poorly on image data? (One sentence)

Ans: An ANN fails because flattening the image into a 1D vector destroys the spatial hierarchies and structural relationships between neighboring pixels

3. Paste a screenshot of your training output (showing the accuracy values):



Step 7: Build a CNN

Now we use Convolutional and Pooling layers.

```
cnn_model = keras.Sequential([
    # First Conv layer: 32 filters, 3x3 kernel
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(
        IMG_HEIGHT, IMG_WIDTH, 3)), layers.MaxPooling2D((2, 2)),
    # Second Conv layer: 64 filters
    layers.Conv2D(64, (3, 3),
        activation='relu'), layers.MaxPooling2D((2, 2)),
    # Third Conv layer (YOUR TURN): Add 64 filters with 3x3 kernel, followed by Max pooling # YOUR
    CODE HERE:
    #
    #
    # Flatten and dense layers
    layers.Flatten(),
    layers.Dense(64, activation='relu'), layers.Dense(1,
        activation='sigmoid') ])
# Compile (copy from ANN) # YOUR
CODE HERE:
#
#
#
#
# Hint: use optimizer = adam, loss = binary_crossentropy, and metrics = accuracy

# Train the CNN
history_cnn = cnn_model.fit(train_generator, validation_data=
    validation_generator, epochs=15)
```


1. What was the final validation accuracy of the CNN?

Answer: 58.16%

2. Compare ANN vs CNN. Which one performed better and why?

Ans: CNN performed better because it uses convolutional filters (kernels) to preserve spatial features and detect local patterns like edges and textures.

3. Paste the accuracy plot comparing ANN and CNN: (Use the code below to generate the plot)



```
plt.plot(history_ann.history['val_accuracy'], label='ANN  
Validation', marker='o') plt.plot(history_cnn.history['val_accuracy'], label='CNN  
Validation', marker='s') plt.xlabel('Epoch')  
plt.ylabel('Validation Accuracy') plt.legend() plt.title('ANN vs CNN  
Validation Accuracy') plt.grid(True) plt.show()
```

Step 8: Fix Overfitting (Most Important)

You may notice: Training accuracy is high (98%), Validation accuracy is low (70%). This is OVERFITTING. The model memorized the training images but doesn't generalize.

Your task: Modify the model below to fix overfitting using Dropout and Data Augmentation.

8.1 Fix 1: Add Dropout

```
from tensorflow.keras import regularizers

# Add dropout of 0.25, 0.25, and 0.5 after each MaxPooling layer respectively
improved_model = keras.Sequential([ layers.Conv2D(32, (3, 3), activation='relu', input_shape=(
    IMG_HEIGHT,      IMG_WIDTH,      3)),
    layers.MaxPooling2D((2, 2)), layers.Dropout( ), #
    Add dropout

    layers.Conv2D(64, (3, 3), activation='relu'), layers.MaxPooling2D((2, 2)),
    #
    , # Add dropout

    layers.Conv2D(64, (3, 3), activation='relu'), layers.MaxPooling2D((2, 2)),

    layers.Flatten(), layers.Dense(64, activation='relu'),
    # , # Add dropout before output layers.Dense(1, activation='sigmoid') ])
# Recompile and Train the improved model (after addressing overfitting using dropout)
improved_model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

history_improved = improved_model.fit(train_generator, validation_data
    =validation_generator, epochs=15)
```

8.2 Fix 2: Add Data Augmentation

Modify the ImageDataGenerator to create more variations:

```

augmented_datagen = ImageDataGenerator(
    rescale=1./255, rotation_range=20,
    width_shift_range=0.1, height_shift_range=0.1,
    horizontal_flip=True, validation_split=0.2
)

train_augmented = augmented_datagen.flow_from_directory( train_path,
    target_size=(IMG_HEIGHT, IMG_WIDTH), batch_size=BATCH_SIZE,
    class_mode='binary', subset='training'
)

# Then retrain the improved model with augmented_datagen
final_model = keras.Sequential([
    # Same architecture as improved_model
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(
        IMG_HEIGHT, IMG_WIDTH, 3)), layers.MaxPooling2D((2, 2)),
    layers.Dropout(0.25), layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)), layers.Dropout(0.25), layers.Conv2D(64,
    (3, 3), activation='relu'), layers.MaxPooling2D((2, 2)), layers.Flatten(),
    layers.Dense(64, activation='relu'), layers.Dropout(0.5),
    layers.Dense(1, activation='sigmoid')
])

final_model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

history_final = final_model.fit(train_augmented, validation_data= validation_generator, epochs=15)

```

1. Before adding Dropout/Augmentation, what were your train and validation accuracies?

.

Train:89.25%Validation 58.16%

2. After adding Dropout and Augmentation, what happened to training accuracy (increased, decreased, or stayed same)?

Train: 60.50%

Training accuracy decreased which is good because it shows the model is no longer memorizing specific images


3. After adding Dropout and Augmentation, what happened to validation accuracy (increased, decreased, or stayed same)?

Validation: 67.35%

Validation accuracy increased

Question 8.3

Paste the before/after validation accuracy plot here:



```
plt.plot(history_cnn.history['val_accuracy'], label='CNN (Overfit)', marker='o')
plt.plot(history_final.history['val_accuracy'], label='CNN + Dropout +
    Augment', marker='s') plt.xlabel('Epoch')
plt.ylabel('Validation Accuracy') plt.legend()
plt.title('Improving Generalization') plt.grid(True)
plt.show()
```

Step 9: Test on Real Images

Find a hot dog image online (or take a photo) and test your model:

```
from tensorflow.keras.preprocessing import image import requests from
io import BytesIO

def predict_image(img_url): # Download image
    response = requests.get(img_url)
    img = image.load_img(BytesIO(response.content), target_size=(
        IMG_HEIGHT, IMG_WIDTH)) img_array = image.img_to_array(img)
    img_array = np.expand_dims(img_array, axis=0) / 255.0

    # Predict prediction = final_model.predict(img_array)[0][0] if prediction >
    0.5:
        print(f" HOT DOG! (confidence: {prediction:.2f})")
    else:
        print(f" NOT A HOT DOG! (confidence: {1-prediction:.2f})")

    # Show image plt.imshow(img)
    plt.axis('off') plt.show()

# Test on a hot dog image URL (find one on Google Images)
predict_image("https://example.com/hotdog.jpg") # REPLACE WITH REAL URL
```

1. Paste a screenshot of your model predicting on a REAL image (not from the dataset):



2. Did your model ever make a funny mistake (e.g., calling a hamburger a hot dog)? What did it confuse?

Ans: My model initially confused a taco with a hot dog. It struggled with items that have similar shapes or bread-like textures but lack the specific visual patterns of a sausage in a bun

Step 10: Optional Stretch Goal YOLO Object Detection (Extra Credit)

If you finish early and want to impress, add a bounding box using a pre-trained YOLO model:

```
!pip install ultralytics -q from ultralytics import YOLO

# Load a tiny pre-trained model yolo_model = YOLO("yolov8n.pt") # n = nano (fast, small)


# Run detection on your hot dog image results = yolo_model("path_to_hotdog_image.jpg")

# Show results (it will draw boxes around objects it recognizes) results[0].show()
```

Extra Credit Question

Does YOLO correctly detect a "hot dog" as a class? What label does it give?

Ans: Yes, YOLO correctly gives the label "hot dog" with a confidence score of **0.76**.



Final Submission Checklist

Step 2: TensorFlow version answer

Step 3: Image counts and balance explanation

Step 4: Screenshot of 6 images + shape answer

Step 6: ANN accuracy + explanation + screenshot

Step 7: CNN accuracy + comparison + plot screenshot

Step 8: Overfitting before/after numbers + plot screenshot

Step 9: Real image prediction screenshot + mistake explanation

Step 10: (Extra credit) YOLO results screenshot

Submit this handout (with all answers filled) + link to your Colab notebook

Remember

The exam will ask you specific questions about YOUR project.
If you run the code yourself, you will know the answers easily.
If you copy from someone else, you will fail the exam.

Goodluck and have fun!